WayVion — PR/FAQ

PRESS RELEASE

WayVion Launches AI-Powered Learning Platform That Adapts to Every Learner in Real Time

Personalized roadmaps, real resources, and adaptive curriculum — built for how people actually learn

Seattle, WA — March 2026 — Today, WayVion announced the launch of its AI-powered learning companion that generates personalized learning roadmaps for any topic, at any skill level, within any timeframe. Unlike static online courses or generic ChatGPT responses, WayVion creates structured, week-by-week curricula with real-world resources — then continuously adapts them based on the learner's actual progress.

"90% of online learners drop out. People spend a third of their learning time just searching for the right resource. And the $190 billion EdTech market still gives everyone the same course," said the WayVion team. "WayVion fixes this by acting like a personal curriculum designer that watches how you're doing and evolves the plan in real time."

The problem: Self-directed learners waste hours searching for scattered resources, following outdated tutorials, or abandoning overly ambitious study plans. Only 5-15% of MOOC learners complete their courses. Workers spend 1.8 hours per day just searching for the right information (McKinsey). Existing platforms offer rigid, one-size-fits-all courses that don't account for individual schedules, goals, or proficiency levels.

The solution: WayVion asks four simple questions — What do you want to learn? What's your current level? What's your goal? How much time do you have? — then generates a complete roadmap with:

- **Structured modules** sequenced by prerequisites and difficulty
- **Real, validated resources** (videos, articles, documentation) sourced from the web and verified via YouTube and search APIs — not AI-hallucinated links
- **Adaptive quizzes** that assess understanding after each module
- **Automatic plan adjustments** when a learner struggles or excels — adding review modules, compressing easy content, or restructuring the timeline
- **Collaborative learning** through study groups where educators assign roadmaps and track group progress

WayVion supports four distinct learning modes — conceptual understanding, hands-on building, quick overview, and deep mastery — across three proficiency levels, creating 12 unique curriculum strategies per topic.

The impact: Research shows personalized learning improves test scores by 54% and boosts retention by 15-20%. Courses with community support see 30-40% higher completion rates. WayVion combines both — personalized AI paths plus social accountability.

"I set up a 2-week JavaScript roadmap as an intermediate learner going for hands-on practice," said a beta tester. "When I struggled with closures and scored low on the quiz, WayVion automatically added a review module and adjusted my timeline. I didn't have to figure out what to do next — it just knew."

WayVion is available now at wayvion.com. Learners can sign up for free or try the platform instantly as a guest.

COMPETITIVE COMPARISON

Feature	WayVion	ChatGPT / Claude	Coursera / Udemy
Personalized to your level, goal & timeline	✅ **Built around you**	⚠️ Approximate	⚠️ Course-level only
Real, verified resources	✅ **Live search, AI-scored**	❌ Hallucinates links	✅ Their library only
Content from the entire internet	✅ **YouTube, docs, articles**	❌ No	❌ Platform locked
Adapts as you progress	✅ **Restructures in real time**	❌ Resets every session	❌ Fixed syllabus
Remembers where you left off	✅ **Fully persistent**	❌ Conversation dies	✅ Maintains Progress
Time-budgeted to your schedule	✅ **Exact hrs/day math**	⚠️ Vague estimates	❌ Pre-defined course duration
Quizzes & accountability	✅ **Per-module, adaptive**	❌ None	✅ Per-module or End-of-course
Restructures when you struggle	✅ **Inserts review modules**	❌ No	❌ No
Learns your behavior over time	✅ **Feedback loop compounds**	❌ No memory	❌ No
Cost	✅ **One platform, any topic**	Free/$20 per month	$15-$80 per course

---

FREQUENTLY ASKED QUESTIONS

External FAQ (Customer / User Questions)

Q: What is WayVion? A: WayVion is an AI-powered learning companion that creates personalized, structured learning roadmaps for any topic. You tell it what you want to learn, your skill level, your learning goal, and how much time you have — and it builds a complete week-by-week plan with real resources, quizzes, and automatic adaptation.

Q: How is this different from just asking ChatGPT to make me a study plan? A: Three key differences:

1. **Real resources, not hallucinated ones.** WayVion uses a 10-stage pipeline with Serper API for real-time web search and YouTube Data API v3 for validation. AI only scores discovered resources — it never generates URLs.
2. **It adapts.** After each module, WayVion checks your quiz scores and self-reported difficulty, then automatically adjusts your plan — adding review modules if you struggled or compressing content if you're ahead.
3. **It persists.** Your progress is tracked, your roadmap is saved, you can share it, join study groups, and pick up exactly where you left off. ChatGPT conversations are disposable.

Q: What topics can I learn? A: Anything. WayVion works for technology (programming, cybersecurity, cloud computing), business (strategy, marketing, finance), healthcare, design, data science, and more. The AI analyzes the domain before building the curriculum.

Q: What are the four learning modes? A:

- **Conceptual** — Focus on "why" and "how it works." 80% reading/watching, 20% reflection.
- **Hands-on** — Build things. 30% learning, 70% doing.
- **Quick Overview** — Hit key points fast. Minimal modules, no deep dives.
- **Deep Mastery** — Comprehensive coverage including edge cases, advanced patterns, and real-world scenarios.

Q: How does the timeline work? A: Three options:

- Set a number of **weeks** + hours per day (e.g., "4 weeks, 2 hours/day")
- Set a number of **days** + hours per day, with optional weekend exclusion
- Set **total hours only** (e.g., "I have 6 hours") for a single-session roadmap

You can also set a hard deadline, and WayVion will plan conservatively to leave buffer time.

Q: What happens if I fall behind or find something too easy? A: Two adaptation mechanisms:

- **Automatic (check-in):** After completing each module, if you report "hard" and score below 60%, WayVion inserts a focused review module. If you report "easy" and score above 90%, it compresses upcoming content. Resources populate in adapted modules only after you accept the change.

- **Manual (adapt-roadmap):** You can request timeline/hours changes anytime. The Architect agent redesigns remaining modules to fit your new budget, preserving completed work.

Q: Are the resources real? How do I know the links actually work? A: Zero-hallucination design. WayVion does NOT generate resource URLs with AI. It uses Serper API for real-time web search, enriches with YouTube Data API v3 for exact metadata (duration, views, channel), enforces anchor-term presence for technical relevance, and applies authority tiering (MDN, FreeCodeCamp, CS50 rank highest). Each URL appears max once across the entire roadmap.

Q: Can I learn with others? A: Yes. Educators and managers can create **study groups**, generate invite codes, and assign roadmaps to the entire group. Each member gets their own cloned copy with individual progress tracking. Group dashboards show everyone's progress with visual rings, average quiz scores, and module difficulty insights. Creators can also invite their WayVion connections directly.

Q: Is it free? A: WayVion offers free access including guest login to try the platform without creating an account.

Q: Can I use this on mobile? A: WayVion is a responsive web application that works on mobile browsers. There is no native app yet.

Internal FAQ (Stakeholder / Technical Questions)

Q: What is the core technical architecture? A: React 18 + TypeScript + Vite frontend, with Supabase providing the entire backend: PostgreSQL database, Auth (email/password + anonymous), Deno-based Edge Functions for AI logic, and Storage. TanStack React Query handles client-side data fetching. shadcn/ui + Tailwind CSS for the design system.

Q: Why Supabase instead of building a custom backend? A: Speed of development. Supabase provides auth, database, edge functions, storage, real-time subscriptions, and Row-Level Security out of the box. For a 4-person team on a university project timeline, this let us focus entirely on the AI-powered features instead of building infrastructure.

Q: What AI models power WayVion? A: 4 distinct AI agent roles distributed across 5 edge functions, using multiple Google Gemini models selected by task complexity:

1. **Agent 1 - Curriculum Architect:** Gemini 3 Pro Preview (roadmap structure, adaptations — most capable, handles 12 goal/level combinations)
2. **Agent 2 - Resource Curator (team of agents):** Gemini 2.5 Flash (resource scoring, reranking — runs as parallel instances, one per module)
3. **Agent 3 - Performance Evaluator:** Gemini 2.5 Flash (check-ins, adaptation triggers)
4. **Agent 4 - Quiz Generator:** Gemini 2.5 Flash Lite (on-demand assessments — lightweight, 3-5 MCQ per module)

The create-roadmap and adapt-roadmap functions are composite services that reuse both the Architect and Curator roles. A 7-module roadmap fires 7 parallel AI calls for resource scoring instead of processing sequentially.

We use dual routing: direct Gemini API first (faster), falling back to the Lovable AI gateway if that fails.

Q: How does roadmap generation work end-to-end? A: Two-phase pipeline:

1. **Phase 1 — Structure (5-15 sec):** The `generate-roadmap` edge function sends a priority-based system prompt to Gemini. The prompt encodes learning goal, skill level, time constraints, and a module count formula. The AI returns structured JSON. Post-processing normalizes module counts, redistributes day ranges, and caps hours per time window. The roadmap is returned immediately with `resources_pending: true`.
2. **Phase 2 — Resources (background):** The `populate-resources` function runs a **10-stage deterministic + AI pipeline**: module context extraction → Serper search queries using anchor terms → YouTube API enrichment → anchor term filtering (relevance gate) → authority scoring → AI scoring (pedagogical fit, level alignment) → diversity checks → span negotiation → per-module parallel AI selection → coverage repair.

This split means users see their roadmap in seconds instead of waiting 30+ seconds.

Q: How do you prevent AI hallucination? A: Resources are never generated by AI. They're discovered via real web search (Serper API), validated via YouTube Data API, filtered through anchor-term gates and allow/exclude lists, and scored for relevance across 6+ dimensions (authority tier, view count, pedagogical fit, level alignment, goal fit, quality signals). The `sanitizeRoadmapPlaceholders` function strips any "search Google for..." text. If JSON parsing fails entirely, a fallback roadmap is generated algorithmically. Semantic deduplication with 0.92 similarity threshold prevents near-duplicate resources.

Q: How does adaptation work? A: Two mechanisms:

1. **Automatic (check-in):** After each module completion, rules-based triggers evaluate self-report + quiz score. "Hard + quiz < 60%" inserts a review module. "Easy + quiz > 90%" compresses upcoming content. Resources populate in adapted modules only after the user accepts the adaptation.
2. **Manual (adapt-roadmap):** User-initiated restructure. The Architect receives completed modules (preserved as-is) and generates entirely new modules for the remaining time budget. Curator fetches new resources for adapted modules. Users can change topic, skill level, timeline, or hours per day mid-roadmap. Both types respect the user's Learning Goal to maintain style consistency.

Both mechanisms recalculate timelines from `total_hours / hours_per_day` — never arbitrary "add a week" adjustments.

Q: How is data security handled? A: Row-Level Security (RLS) on every Supabase table. Users can only access their own roadmaps, progress, and group data via `auth.uid()` policies. Group policies scope visibility to members only. The Supabase anon key is public by design — all data access is controlled by RLS. Server-side secrets (API keys) are only accessible in edge functions via `Deno.env.get()`.

Q: How do you prevent abuse? A: Three layers:

- **Rate limiting:** In-memory sliding window per user per function (10/hr for roadmap generation, 20/hr for adaptation, 30/hr for quizzes, 60/hr for check-ins).
- **Input sanitization:** All user inputs truncated to max length and stripped of dangerous characters before being injected into prompts.
- **Auth enforcement:** Every edge function verifies the JWT token before processing.

Q: What were the hardest engineering challenges? A:

1. **Zero-hallucination resource curation (~1,600 lines):** Engineering the 10-stage pipeline to prevent AI from generating fake URLs while maintaining quality required multiple fallback layers, anchor-term gates, mathematical time budgeting, and parallel AI scoring across modules.
2. **Token budget management:** Keeping Gemini context under limits during check-in while passing enough roadmap data for intelligent decisions. Required aggressive pruning of completed modules and resource metadata.
3. **Group isolation vs. updates dilemma:** Chose full roadmap cloning for member safety, but this means owners can't push updates to already-shared roadmaps (still unsolved).
4. **RLS policy correctness for groups:** 5+ rounds of fixes to get group member visibility, admin permissions, and invite-based joining all working correctly without creating security holes.
5. **AI output unpredictability:** Even with structured JSON mode, the AI sometimes returns malformed output. Built multi-attempt parsing (raw → fence extraction → control char sanitization → combined) plus algorithmic fallback generation.

Q: What trade-offs did you make? A:

1. **JSON Blob vs. Normalized Schema —** Roadmap data stored as JSON blob. Gain: AI can generate arbitrary structures without migrations. Cost: Can't query by module properties.
2. **Roadmap Cloning vs. Shared Instances —** Full copies for each group member. Gain: Members adapt independently. Cost: Storage bloat; no canonical version.
3. **Deferred Resource Population —** Show roadmap first, populate resources in background. Gain: <3s initial response. Cost: Two-phase UX.
4. **AI Model Tiering —** 3 different Gemini models. Gain: 60-70% cost reduction on quizzes. Cost: Complex fallback logic.
5. **Hardcoded Authority Tiers —** MDN=0.95, FreeCodeCamp=0.9. Gain: Predictable quality. Cost: Doesn't adapt to emerging sources.
6. **Single Search Provider (Serper) —** One API. Gain: Consistent format. Cost: Single point of failure.
7. **On-Demand Quiz Generation —** Created at completion time. Gain: Fresh questions. Cost: 2-5s delay.

Biggest architectural bet: Flexibility over performance — JSON storage, cloned roadmaps, and deferred curation prioritize rapid AI iteration at the expense of query efficiency and storage costs.

Q: How would this scale beyond a university project? A:

- Move rate limiting from in-memory to Redis for cross-isolate accuracy
- Add a job queue (e.g., Supabase pg_cron or external queue) for background resource population
- Cache popular roadmap templates for instant delivery on common topics
- Split the resource pipeline into a dedicated service
- Add OAuth providers (Google, GitHub) beyond email/password
- Build native mobile apps

Q: People come for the roadmap. What makes them stay? A: Community. WayVion already has the social infrastructure — profiles, connections, groups, progress tracking, roadmap sharing. Courses with community support see 30-40% higher completion rates. Our groups feature gives educators real-time visibility into who's struggling and which modules are too hard, creating accountability loops that keep learners engaged.

What's next is the engagement depth layer: discussion threads per module/group (peer Q&A), a roadmap marketplace (discover what others built), public resource reviews (community-driven quality signals), activity feeds, leaderboards, and badges. The foundation — database schema, RLS policies, role system — is designed to support all of this. We built the moat (AI pipeline) first; now we're building the city (community).

Q: What's on the future roadmap? A:

Community & Engagement (next priority):

1. **Discussion/Q&A system —** Per-module and per-group threads for peer support within the learning flow
2. **Roadmap marketplace —** Browse, discover, and clone top-rated community roadmaps
3. **Public resource reviews —** Aggregate community ratings on resources ("87% found this helpful")
4. **Activity feed —** See what connections are learning and completing
5. **Leaderboard & badges —** Weekly rankings, milestone achievements

AI & Content Quality: 6. **Transcript-Based Scoring —** Pull YouTube transcripts and article content to score against learning objectives. Currently judging resources by wrapper (title, metadata) not content. 7. **"Why This Resource?" Explainer —** AI-generated one-liner on each resource card explaining why it was selected, making curation transparent.

Retention & Growth: 8. **Learning Streak + Daily Digest Email —** Push notifications tied to specific roadmaps and deadlines to combat abandonment. 9. Spaced repetition improvements for flashcards 10. Richer analytics dashboard for educators 11. OAuth providers (Google, GitHub) 12. Landing page overhaul with AI-generated topic images

Q: What's the team structure? A:

Role	Person	Responsibility
CEO / Founder	Swati	Architecture, core development, live demo
COO	Sandrine	Operations, security, architecture Q&A
CFO	Nick	Financial planning, AI features Q&A
VP Marketing & Sales	Nathan	Project presentation, vision, go-to-market

Q: What metrics define success? A:

- **Engagement:** % of users who complete at least 1 module after generating a roadmap
- **Adaptation effectiveness:** Quiz score improvement after adaptation triggers
- **Resource quality:** User feedback ratings on suggested resources; % of dead/irrelevant links
- **Retention:** % of users who return to continue a roadmap within 7 days
- **Group adoption:** Number of groups created, average group size, group completion rates